

CS 208 - Applied Privacy for Data Science Homework 2

Jason Huang
Spring 2019 - Harvard University

The public Github repo containing all work is at <https://github.com/TurboFreeze/cs208hw>. All code has also been included in the appendix of this PDF as specified.

Problem 1

- (i) (a) The clamping function is effectively applying a post-processing function to the noisy query result. In other words, Laplace noise is added to the true mean \bar{x} , which must be $(\epsilon, 0)$ -DP. The following clamping function does not change the privacy characteristics guaranteed by differential privacy, meaning that this mechanism **meets the definition** of $(\epsilon, 0)$ -DP (following directly by privacy under post-processing and the proof of Laplace DP).

Note that the scale factor parameter of the Laplace distribution should be set to $s = GS_q/\epsilon$ for differential privacy, meaning that $\epsilon = GS_q/s$. In this case, the global sensitivity GS_q is the maximum change that can be affected to the statistic by a single entry's change, which in this case would be $1/n$ for the mean. Furthermore $s = 2/n$. The $\epsilon = (1/n)/(2/n) \implies \boxed{\epsilon = 0.5}$.

- (ii) Constant ratios of Laplace mechanisms

- (iii)

$$\frac{P[M(x', q) = r]}{P[M(x, q) = r]} =$$

- (iv)

$$\begin{aligned} P[M(x, q) = r] &= P[\bar{x} + Z]_0^1 = r] \\ &= \\ \frac{P[M(x, q) = r]}{P[M(x', q) = r]} &= \\ P[M(x, q) = r] &= P[\bar{x} + [Z]_{-1}^1 = r] \\ &= \end{aligned}$$

- (d)

Problem 2

(a) The DGP is the following likelihood of some data vector $k \in \mathbb{N}^n$:

$$P(\mathbf{x} = \mathbf{k}) = \prod_{i=1}^n \frac{10^{\mathbf{k}_i} e^{-10}}{\mathbf{k}_i!}$$


The DGP function was implemented using a Poisson random draw.

See the attached R script `q2.R` for the implementation.

(b) The first mechanism was chosen, involving clamping after Laplace noise has been added.

See the attached R script `q2.R` for the implementation.

(c) The optimal value b^* for b is $\boxed{b^* \approx 15}$. As expected, root mean squared error is indeed high with small clamping regions and decreases as it becomes more appropriate, with large clamping regions yielding high RMSE again.



See the attached R script `q2.R` for the implementation.

(d) This approach is not safe and might violate differential privacy because it implicitly incorporates information about the dataset that is not included in the data. More specifically, bootstrapping is powerful because of its very ability to simulate the data generating process and the distribution of the data. Therefore, bootstrapping to find an estimate for the optimal value b^* is effectively the same thing as calculating the true optimal b^* , only numerically instead of analytically. Now, knowing the optimal b^* violates privacy because it also provides insight into the distribution (i.e. one of the ways this can be intuitively described is as an approximate maximum of everything that is not an outlier). This is similar to the idea of why local sensitivity, which is a characteristic of the data set, is not

(ϵ, δ) -DP but global sensitivity is. The optimal b^* is a specific data-dependent attribute of the data set that can leak privacy if not controlled for, as is the case with bootstrapping here.

(e) Use an external reference for determining an approximation for b^* . Generally, a similar public data set (i.e. past iterations of the census or different geographic tracts that have been made public) can help provide a rough estimate of the upper bound for the individual. Often times, a reasonable heuristic can be used based on common knowledge, such as the rough range of human ages or other metadata provided for the data set (i.e. in a codebook or schema).

Problem 3

(a) There are differentially private techniques to release the means \bar{y} and \bar{x} as well as the slope $\hat{\beta}$. However, given the careful considerations needed for the slope $\hat{\beta}$, it may be challenging to come up with a single differentially private mechanism to derive the intercept estimate. However $\hat{\alpha} = \bar{y} - \hat{\beta}\bar{x}$ lends itself nicely to privacy preservation under composition and post-processing. Since there are three differentially private statistics needed here of $\bar{x}, \bar{y}, \hat{\beta}$, for a total epsilon budget of ϵ_t , then calculate differentially private releases of each statistic with $\epsilon = \epsilon_t/4$ (note that the slope $\hat{\beta}$ is actually two statistics of covariance and variance in the numerator and denominator respectively) to yield $(\epsilon_t/4, 0)$ -DP statistics. Since post-processing is allowed without affecting privacy, then this three differentially private statistics will lead to $\epsilon_t/4 + \epsilon_t/4 + \epsilon_t/2 = \epsilon_t$ differential privacy for $\hat{\alpha}$ by composition and $\epsilon_t/4 + \epsilon_t/4 = \epsilon_t/2$ differential privacy for $\hat{\beta}$ (which is used in $\hat{\alpha}$ and does not require separate consumption of the budget). Therefore, the overall method for computing both $\hat{\alpha}$ and $\hat{\beta}$ would be ϵ_t -DP, as desired.

Since the data x_i is generated by a Poisson process according to the previous problem, it can be clamped using the optimal value of $b^* \approx 10$ found before. Since there is a linear relationship between x_i and y_i here (and it is known in the following part that the slope is simply 1), then similarly clamp y_i by $b^* \approx 15$.

See the attached R script `q3.R` for the implementation.

(b)

See the attached R script `q3.R` for the implementation.

(c)

See the attached R script `q3.R` for the implementation.

Problem 4

Use linearity of expectations and fundamental bridge to convert between probabilities and expectation of indicators.

$$\begin{aligned} \mathbb{E}[\#\{i \in [n] : A(M(X))_i = X_i\}/n] &= \mathbb{E}[\mathbb{1}\{i \in [n] : A(M(X))_i = X_i\}/n] \\ &= \mathbb{E}\left[\sum_{i=1}^n \mathbb{1}(A(M(X))_i = X_i)/n\right] \end{aligned}$$

$$\begin{aligned}
 &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}[\mathbb{1}(A(M(X))_i = X_i)] \\
 &= \frac{1}{n} \sum_{i=1}^n P(A(M(X))_i = X_i) \\
 &= P(A(M(X))_i = X_i)
 \end{aligned}$$

Where the last simplification is made by noting the independence of the X_i data draws. Let there be a dataset X^{i0} or X^{i1} such that the i^{th} row is changed to 0 or 1 respectively, with corresponding probabilities of p and $1 - p$ by the nature of the given Bernoulli data generating process. Now by the definition of (ϵ, δ) -DP:

$$\begin{aligned}
 \mathbb{E}[\#\{i \in [n] : A(M(X))_i = X_i\}/n] &= P(A(M(X))_i = X_i) \\
 &\leq e^\epsilon \cdot \max\{P(A(M(X^{i0}))_i = X_i), P(A(M(X^{i1}))_i = X_i)\} + \delta \\
 &\leq e^\epsilon \cdot \max\{p, 1 - p\} + \delta
 \end{aligned}$$

The desired result has been shown.

$$\therefore \boxed{\mathbb{E}[\#\{i \in [n] : A(M(X))_i = X_i\}/n] \leq e^\epsilon \cdot \max\{p, 1 - p\} + \delta} \text{ as desired}$$

Appendix

Code for Problem 1

Code for Problem 2

```
##  
## q2.r  
##  
## Evaluating DP algorithms with synthetic data  
##  
## JH 2019/03/10  
##  
  
##3 PART (A)  
# data generating process  
dgp <- function (n, lambda=10) {  
  rpois(n, lambda)  
}  
  
#### PART (B)  
  
# sign function  
sgn <- function(x) {  
  return(ifelse(x < 0, -1, 1))  
}  
  
# laplace random draws  
rlap <- function(mu=0, s=1, size=1) {  
  p <- runif(size) - 0.5  
  draws <- mu - s * sgn(p) * log(1 - 2 * abs(p))  
  draws  
}  
  
# clamping helper function  
clamp <- function (data, a, b) {  
  data.clamped <- data  
  data.clamped[data < a] <- a  
  data.clamped[data > b] <- b  
  data.clamped  
}  
  
# differentially private mechanism (clamping)  
dpclamping <- function (data, epsilon, a=0, b) {  
  mean.actual <- mean(clamp(data, a, b))
```

```
# generate noise by Laplace mechanism
data.len <- length(data)
sensitivity <- (b - a) / data.len
laplace.shift <- sensitivity / epsilon
noise <- rlap(s=laplace.shift, size=1)

# inject noise
mean.noisy <- mean.actual + noise

# apply clamping
mean.clamped <- clamp(mean.noisy, a, b)
}

#### PART (C)

# parameters
n <- 200
epsilon <- 0.5
b.seq <- seq(0, 30, by=0.1)
b.num <- length(b.seq)

# generate data
data.poisson <- dgp(n, lambda=10)

# calculate RMSE
rmse <- function (data, m) {
  sqrt(sum((data - m)^2))
}

n.trials <- 100
rmse <- vector("numeric", b.num)
for (i in 1:b.num) {

  dpmeans <- vector("numeric", n.trials)
  for (j in 1:n.trials) {
    # calculate dp query of mean
    dpmeans[j] <- dpclamping(data=data.poisson, epsilon=epsilon, b=b.seq[i])
  }
  # calculate and store RMSE
  rmse[i] <- rmse(dpmeans, m=mean(data.poisson))
}
```

```
# find index of minimum RMSE (optimal value of b)
b.optimal <- b.seq[which.min(rmse)]; b.optimal

#plot(b.seq, rmse, type='l') # uncomment if no ggplot

# visualize
rmsedata <- data.frame(b=b.seq, rmse=rmse)
library(ggplot2)
clamping.plot <- ggplot(rmsedata, aes(x=b, y=rmse)) +
  geom_line() +
  geom_vline(xintercept=b.optimal, color="red", alpha=0.8, linetype="dashed")
  theme_bw()
clamping.plot
ggsave("clamping.jpg", clamping.plot)
```


Code for Problem 3